

Lecture 2:
Intro to CPU
Mechanisms

Virtualization: The CPU [Fall '11] ①

goal: run N processes @ once,
even though only m CPUs
($N \gg m$)

[multi-programming]

CPU: A | B | C | A | B | C
time sharing


⇒ also, it's own

restricted
machine

⇒ it's own
efficient
machine

abstraction: the Process (running program)

Q what does a running prog do? private memory: address space
(kind of ignore for now)



virtualizing CPU
⇒ virtualizing registers
of machine

what is OS to a process?
⇒ set of APIs
(std library)

to support abstraction: need

focus: mechanisms how to virtualize policies which P to run?
(next time)

general mech: limited direct execution

How to virtualize CPU:

start simple: when you want to run a program, just run it!
(load into memory, point PC @ first inst, go!)

⇒ direct execution

SHOW
PROTOCOL
HERE
(LT)

Three Problems:

- 1) what if process wants to do something restricted?
- 2) what if OS wants to stop P_A , run P_B ?
- 3) what if running P does something slow, like I/O?

Problem #1: restricted ops (e.g. I/O)


solution: privilege levels + way to change it
(e.g. the system call)

P_A P_B
↓ ↓
[F_A] [F_B]
to read:
is instruction
(read block 10)

could trust processes
not to read other user's
files....

1) user process: "user" mode (restricted) 2) OS: "kernel" mode (not restricted)

need way to change between them: sys call



⇒ save state of running process
trap: changes privilege
⇒ level to kernel mode
⇒ runs "trap handler"
⇒ code for handling
this trap
rett: ⇒ restore state
priv. level run back...

Prob #2: cont (Fall '11, ②)

super
powerful
instruction!

- 1) save state: done by h/w (+ s/w) ^{sometimes}
logically: save PC, registers to some special place
then, can run OS code w/o worry of
corrupting user code registers
- 2) raise priv. level: allows fun stuff
- 3) jump to correct kernel code: how does h/w know
which PC to jump to?

Q some save state
must be done
by h/w. why?

Q special instruction:
privileged?

protocol:

kernel boots (kernel mode)

installs trap tables ~~list of addresses~~
of functions to run when certain
traps/interrupts happen

eventually runs a process (user mode)

sys calls: → kernel on trap:
user ← h/w consults trap
table to know
which code to run

instruction:

ret-from-trap

- 1) lower priv. level
- 2) restore state of registers of calling process
- 3) jump back to PC after trap in user code

OS issue:
where to
put saved
state?
⇒

Q what if user-mode process tries to do something restricted?
⇒ (killed)

Problem #2: how to stop currently running P_A, run P_B?
(only 1 CPU: OS is not running, when P_A is running)

approach #1: cooperative just wait for process to
call into kernel or yield explicitly

#2: non-cooperative forcibly reclaim CPU via
timer interrupt

issue:
must

it switch: ⇒ context switch
save P_A state, stop it

restore P_B's state, run it

[turn on @ boot, before running
any user process]


goes off every X ms:
decide to continue P_A,
or switch to P_B?

Prob #2: cont [Fall '11, ③]

OS issues:

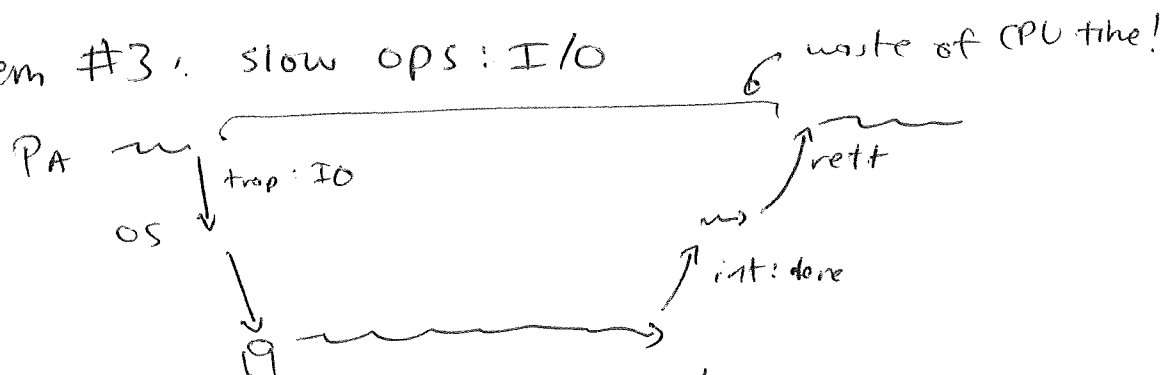
1) trust

2) have to track state of processes




3) need data struct to track this: process list
per-process info: saved registers, state: ready, running

Problem #3: slow ops: I/O



new process state: blocked



when trap into kernel for I/O, OS switches to another process ⇒ efficiency

Summary: Limited Direct Execution

direct: because usually, just runs

limited: at certain key points, OS + h/w get in way to retain control of system

result: ~~small~~ virtualized CPU

⇒ w/ safety

⇒ w/ efficiency