

Lecture 3 :

Scheduling

Last time: CPU virtualization [Fall '11]

Mechanisms: Limited Direct Execution

protocol {  
@ boot: setup trap handlers  
then run jobs:  
⇒ OS involved on sys calls (notion of privilege)  
⇒ OS involved on timer interrupts (context switch)  
result: efficient, but w/ control  
[OS as manager]


Today: Policies (easier to understand)

based on notion of time sharing  
~~notion of time sharing~~ (job ~ running program)  
Assume: simplistic

- ⇒ set of jobs, all arrive @ once
- ⇒ each just uses CPU (no I/O)
- ⇒ each runs for same length of time
- ⇒ time is known & priority
- ⇒ one metric: turnaround time  
def:  $(T_{\text{complete}} - T_{\text{arrive}})$


Algorithm #1: FIFO / FCFS (queue)

A, B, C get queued up; each run 10 time units



But, Relax #3: jobs of different lengths

Q) how does FIFO do? → convoy: C stuck behind B (real life)



Algorithm #2: SJF ⇒  $T_{\text{avg}} = 50$   
optimal (given our assumptions)

But, Relax #2: not all at same time ⇒ problem?

[A: 10, B: 100 arrive @ T=0; C: 10 arrives @ T=20]

Algorithm #3: STCF (allows pre-emption)

But, Relax #5: other metrics: Response Time =  $(T_{\text{first run}} - T_{\text{arrive}})$

Algorithm #4: RR (concept (time slice))

timer interrupt every X ms;  
time slice is n \* X (n ≥ 1)



$R_{\text{avg}} = ?$


$T_{\text{avg}} = ?$  (just intuition)

More Real Workload:  
mix {  
interactive: short CPU bursts, then long pauses (blocked)  
batch: long running CPU workloads - intense

more Real Scheduler:  
can't assume knowledge of jobs  
⇒ what to do? HISTORY

MLFQ: basic idea:  
assume new job is interactive,  
learn if this is true/not over time!

start w/ RR, but many queues:



each job has a priority @ a given time (may change)

Base Rules:

- 1) if  $R(A) > P(B) \Rightarrow A$  runs
- 2) if  $P(A) == P(B) \Rightarrow RR(A, B)$

Key: pri changes over time: how?

3) start jobs @ highest priority

Attempt #1:

4) ~~not~~ if I/O issued before a) time slice is up ⇒ (e.g. I/O) stop @ same pri

b) if not, move down one level

(EXAMPLES)

long running batch, late arriving interactive

Problems? 1) starvation, job classes, 2) Gene the scheduler

Attempt #2: solve starvation

Attempt #3: solve gaming

(Also: time slice length)

But: what about I/O? (interactive jobs use short bursts of CPU)