

LECTURE 8:

The Page Fault
(+ Policy)

Last time(s) :

Paging:

too slow \rightarrow TLB

too big \rightarrow multi-level PT

Today:

what if all that stuff
doesn't fit into memory?

Ⓚ when does this occur?


A: one program w/ lots of data
many programs sharing mem

Overview:

{ Mechanisms
Policies }

Mechanisms

Swap Space : + disk




size Ⓚ how big is a modern hard drive?

Thus, page may be in-memory
or on disk

\Rightarrow OS needs to track this for
all pages in system

Ⓚ how? \Rightarrow in page table!

Page Table Entry



new bit:

present bit

if page is in memory \Rightarrow P=1
else P=0


where in phys
this page is

note: when present=0,
can use PFN field

Ⓚ what to use
it for?

New Control Flow

Trying to fetch instruction @ VA



1) extract vPN

2) check TLB for vPN

2a) ^{TLB} hit: get PFN (offset \Rightarrow Fetch PA from memory)

2b) ^{TLB} miss: calculate address of PTE

$PA_{PTE} = (PTBR + (vPN * \text{sizeof}(PTE)))$

\Rightarrow Fetch PA_{PTE} raise

\Rightarrow if not valid \Rightarrow Exception (SEG FAULT)

if prot not OK \Rightarrow " (PROT FAULT)

if present:

extract PFN
insert \Rightarrow TLB
retry

else

Raise Exception (PAGE FAULT)

should be
PAGE MISS

\Rightarrow Page Fault Handler

PFN = Find Free Phys Page()

if (PFN == -1)

my sleep

\Rightarrow Evict Page()

using page
replacement
policy

DiskRead (PTE.DiskAddr,
PFN)

update PTE: (present=1, PFN)

retry instruction

Policy: which page to evict? (replace)

Ⓚ Optimal: kick out page ref'd furthest in future

Random: pick any one, kick it out


Ⓚ FIFO: keep FIFO list of pages
(also, LFU)

Ⓚ LRU: kick out least-recently-used page
note: difference in updates
(LRU updates list on every access)

[DO EXAMPLE]

Conclude: Good to do LRU-like policy
But: how to implement?

Approximate: h/w support
via use bit (reference bit)




upon every access to a page,
h/w sets use bit (to 1)

thus, after running for a while,
can differentiate which
are in use, which aren't

Ⓚ is this enough for whole policy?
no, OS must clear bits (use=0)
periodically too

clock algorithm



clock hand

if use=1
clear (use=0)
inc clock
else
replace (this page)

Finally:

Thrashing: (working set): actively used
pages
memory is over committed
solutions: admission control, kill some
procs

\Rightarrow Buy more memory

inst Ⓚ (does this always
work?)

Belady's Anomaly:

1, 2, 3, 4, 1, 2, 5, 1, 2, 3, 4, 5

w/ FIFO₃ and FIFO₄