

Concurrency #1

- intro to problem
- simple lock

INTRO

CONCURRENCY

Introduction: What is a thread?

- on board
 - multiple "programs" executing within the SAME address space!
 - usually cooperating to achieve some task, or independent related tasks
- e.g., parallel program, web/db server
 - thus, each thread has:
 - its own private set of registers
 - its own program counter
 - its own stack (and sp/bp)
 - but shares
 - rest of address space (heap, static global, code section too)

ATOMICITY

What makes thread programming hard?

main-thread-0 (no locks)

objdump -d main (inspect code)

use this to VISUALIZE address space

(per-thread stacks, shared parts)

Examine in detail: main-trace.txt

Why programs get tricky: SHARED DATA

REAL PROBLEM: uncontrolled scheduling (interrupts at any time)

lots of definitions:

- program is not deterministic (indeterminate)
- critical section
- race condition
- need mutual exclusion (turn indeterminate code into

deterministic code)

main-thread-1 (fine-grained locks)

need synch primitives

main-thread-2 (coarse-grained locks)

need synch primitives but be careful

main-thread-3 (implement locks try #1: test-and-set)

just run it

(what is the problem?)

main-thread-4 (implement locks try #2: x86 xchg)

how to build a lock using special hardware?

(how to use xchg?)

main-thread-5 (implement locks try #2: x86 xchg + spinlock implementation)

this is how

objdump -d to look at it

Conclusions:

Why in OS class?

threads are basic OS primitive AND OS is a concurrent program!