

Concurrence #6:

→ Finish sema (P/C)  
problem

→ Deadlock

```

int max;
int loops;
int *buffer;

int use = 0;
int fill = 0;

#define CMAX (10)
int consumers = 1;

void do_fill(int value)
{
    buffer[fill] = value;
    fill++;
    if (fill == max)
        fill = 0;
}

int do_get()
{
    int tmp = buffer[use];
    use++;
    if (use == max)
        use = 0;
    return tmp;
}


void *
producer(void *arg)
{
    int i;
    for (i = 0; i < loops; i++) {
        do_fill(i);
    }
}

void *
consumer(void *arg)
{
    int tmp = 0;
    while (tmp != -1) {
        tmp = do_get();
        printf("%d\n", tmp);
    }
}


```

# Deadlock

everyone waiting for something held by other  
none release until they get what they're waiting for



e.g. stop sign



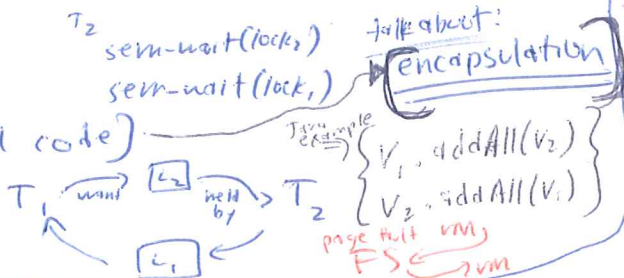
int Kansas: "when two trains approach each other @ crossing, both shall stop and neither start til other has gone"

⇒ code example:

$T_1$  sem-wait(lock<sub>1</sub>)  
sem-wait(lock<sub>2</sub>)

(happens in real code)

$T_1 \rightarrow T_2$  or



talk about: **encapsulation**

example:  
 $V_1 \cdot addAll(V_2)$   
 $V_2 \cdot addAll(V_1)$   
page fault vm  
FS ← vm

## 3) Detect/Recover

⇒ detect: keep track of units-for graph  
if cycle detected,  
→ tell human?  
→ restart thread/system

(DBMS)

challenge: not leaving system in funny state

lock 1  
x subtracted from acct 1  
→ lock 2  
add x to acct 2  
unlock 2  
unlock 1

## Conditions

mutual exclusion: ≥ 1 resource held in non-sharable mode  
another req → delay

hold-and-wait: ∃ process holding, waiting *may request new*

no preemption: can't be preempted

circular wait: ∃  $P_0, \dots, P_n$  s.t.  $P_0 \rightarrow P_1 \rightarrow \dots \rightarrow P_n \rightarrow P_0$

## Solutions

prevent: ensure 1/more conditions doesn't hold  
→ avoid ⇒ scheduling  
→ detect/recover: allow but notice + recover  
→ ignore: not everything worth doing is worth doing well

## Prevent

(maybe read-only?)  
→ mutex: hard to get around this → lock-free structs  
→ hold+wait: guarantee all locks grabbed @ once, atomically  
(but encapsulation makes this hard)  
→ also, starvation possible?

→ no preemption:  
mutex-try/lock() {  
    acquire sem-wait(lock)  
    = sem-pul(lock)  
}

→ circular wait:  
impose total order  
always grab lock<sub>1</sub> before lock<sub>2</sub>  
(common approach)

{ write 2-lock acquire that works even if done in wrong order } ⇒ (how to do w/ encapsulation?)

## wait-free synch.

insert-at-head (int)  
node-t \*new = malloc();  
new → element = element;  
new → next = head;  
head = new;

⇒ is this a critical section?  
(example)

```

int CAS(unsigned *addr, unsigned eap, unsigned val)
{
    if (*addr == eap) {
        *addr = val;
        return 1;
    }
    return 0;
}

```

⇒ how to rewrite list insert s.t. it never uses CAS to update list w/o locks?

```

node-t *new = malloc();
do {
    node-t *old = head;
    new → next = old;
    while (CAS(&old, old, new) == 0);
}

```

negatives:

## Avoidance:

w/ smart scheduler  
if you knew something about maximal possible requests of a thread, could decide how to schedule (Banker's alg, others)

