

LECTURE

FS : Files + Directories

Basic Abstraction: FILE (regular file)

- array of bytes, has low-level name (i.e., a number), can create/read/write/grow/delete, persistent

Basic Abstraction: DIRECTORY (special type of file)

- array of records, maps human-understandable file/dir name to low-level file/dir name, can create/delete/read

**** FILE OPERATIONS ****

CREATION - creat.c - look at it, run it, strace it

ACCESS (read) -

concept of **FILE DESCRIPTOR** - open files are tracked (in table)
(not shared with other processes, in general)

file descriptor: per-process int to refer to given open file
tracked per open file: **CURRENT OFFSET** (inc'd after read/write)
cat file1 -- strace cat file1

make file bigger w/ "echo a >> file1" in loop 5000x **(HOW BIG?)**
strace cat file1 > /dev/null (make sure to TRACK CURRENT OFFSET)

GROW (write) [FILE grow.c]

how do you grow a file? lseek() to seek to end of the file

TRUNCATE (make file smaller) [FILE trunc.c]

strace trunc, ls -al file1, truncate() syscall (not always zero)

DELETE - **ASK how to delete a file?** which system call? MYSTERY!

strace rm file1, show that it is "unlink"! (WHY??)

RENAME - **ASK how to rename a file?** -- mv file1 file2

strace it, see that rename() system call is used

STAT - **ASK how to get info about file?** - strace stat file2

which leads to:: **ASK what info should FS store about a file?**

all information is stored in per-file structure called an **inode**
KEY: low-level name of file is thus the "inode number"

(short for **index node**, because original unix FS inodes in array)

Some other interesting things in there:

SIZE (and BLOCKS), OWNERSHIP, PERMISSIONS, ACCESS TIMES,

POINTERS to BLOCKS (we can't see these)

**** DIRECTORY OPERATIONS ****

Directory: just a **SPECIAL TYPE of FILE**

can create/read/delete (but NOT write to, directly)

stores mapping human-readable name->low-level name (inode#)

CREATE/REMOVE - strace mkdir foo, strace rmdir foo

REaddir - can READ as well - readdir.c

(show how to build this using opendir(), readdir(), closedir())

ASK: what program is this?? (answer: it's "ls")

Some **SPECIAL** directories:

ROOT: "/" at top of tree - **CURRENT:** "." - **PARENT:** ".."

PATHNAMES

ABSOLUTE: /x/y/z -- root / dir / dir / ... / file or dir

RELATIVE: foo/bar.c -- dir / dir / ... / file or dir

key notion: **SEPARATOR** (/ is the right one, others used)

each process has **CURRENT WORKING DIRECTORY** (getcwd())

WELL-KNOWN INODE # FOR ROOT -> ls -ali / -> root is 2

PATH TRAVERSAL - Revisiting OPEN (What does OPEN() do?)

fd = open("/a/b/c/file1", O_RDONLY);

start @ ROOT (or CWD if relative), **TRAVERSE** root directory,

'a', 'b', 'c', finally finds 'file1' and reads its inode

fd is then used to refer to 'file structure' for this file

LINKS (aka HARD LINKS)

NOT a special type of file (but you might think it is)

how to create? -> ln file1 file2 -> strace it! (link())

eg: cat file1 - cat file2 - ls -ali file* -> same inode #

stat file1 file2 -> see LINK COUNT

rm file1 - cat file2 -> **WILL IT WORK?**

stat file2 -> see LINK COUNT

rm file2

And now you know: why REMOVE is done by UNLINKing a file

each unlink removes single reference to file (**REF COUNT--**)

when final unlink occurs, file is removed from file system

LINKS (SYMBOLIC aka SOFT) - a special file type

ln -s file1 file2 -> strace to show symlink() system call

ls -al file1 file2 -> SHOW HOW IT IS A SPECIAL FILE TYPE

NOTE FILE SIZE of file2

(it is 5 bytes: can anyone guess why?)

rm file1 - cat file2 --> dangling pointer!

why needed? (to point to other FSes, dirs in cycle)

**** MOUNTING FILE SYSTEMS ****

each file system is stored in a **"VOLUME"**

looks like a disk (but could be part of a disk, a partition)

but all stitched together into **SINGLE TREE** of all file systems

type **"mount"**

show root directory

show how another disk is mounted on **"/scratch.1"**

can use this to assemble full file-system tree!

(instead of having c: and d: and e: on desktop)