LECTURE

FFS

Review:

Last time: FS implementation

simple unix fs

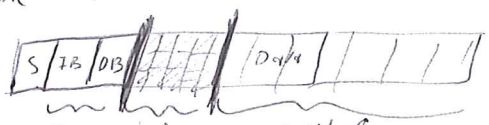

allocation tracking | inodes (per-file/dir metadata) | most of disk partition

today: (bigger files, FFS)

⇒ support for large files, directories?

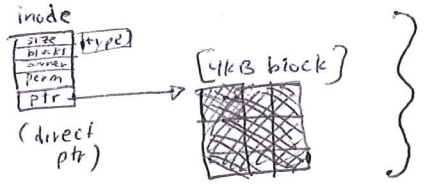

how to make bigger?

- move direct ptrs
- linked (FAT)

- extents (ptr, length)
- indirect blocks (multi-level index)

(example: how to calc block # of offset X)

⇒ FFS fast file system

old unix fs: easy to use but... 2% of disk peak

problem: performance

{keep: same interface
change: underlying impl}

old unix fs:

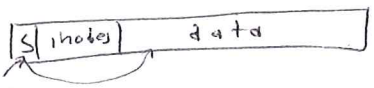

512-byte (small) blocks

problems: (get students to do this)

- → inodes far from data
- → related files: far from each other
- → free list: fragmented (example)

overall: structures are not "disk aware" (treats disk like RAM)

FFS: treat disk like disk!

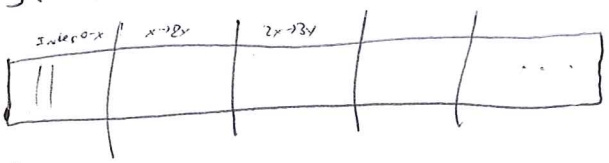

cylinder/block group

1) chop disk into groups
2) put "related" things in same group

(corrolary: "unrelated" in different groups)

Group closeup:

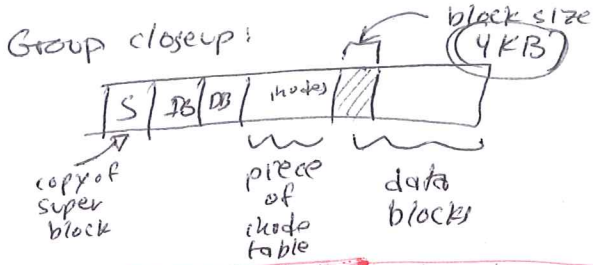

Q) given inode # X, how to calc. on-disk addr?

⇒ Bitmaps not free list: why? (can see if nearby blocks are free easily)

need: allocation policies

what is related?

(Q) always true?

{inode → it's own data
files in same directory
(namespace-based locality)}

mkdir: pick group w/ (high free # of inodes, data blks, low # of directories)

create file:

place data blocks in same group as inode

exception: large files

didn't want to place a large file into one group: why?

(fills up group, can't put future "related" stuff there)

FFS rule: once file reaches certain size, place next chunk of file in some other group:

(how big should chunk be?

{e.g. xfer: 50 MB/s (peak)
avg seek: 10 ms
rot

to achieve 90% of peak b/w for large files?

other FFS fun:

- → sym links
- → long file names
- → parameterization of layout
- → sub-blocks (for really small files) ⇒ avoid internal frag

Summary: treat disk like disk (group related things)