


LECTURE

AFS

AFS (NFS, block based)  
 what is the def?  
 goal: scale  
 (+ cache consistency, + use local disk!)  
 what is the limit of this?  
 Prototype 1: AFS v1



notice:  
 all clients  
 have disks!

FUN IDEA:  
 Reverse engineer  
AFS!

Thus:

open ("/x/y/z.doc")  
 [assume root is local]

For each remote dir  
 "x", "y" =>

1) if in local cache  
 w/ callback,  
use w/o comm

2) if in cache but  
 w/o callback  
 [e.g. after client  
 reboot, server reboot,  
 net failure]

ask server for  
 new callback

3) not in cache,  
get it, get callback

Next open =>  
all local!

Crash Recovery: harder

= {why?}

server crash:  
 what happens to  
 all the state?

{what if just a  
 network failure?}

=> {availability vs,  
 consistency vs,  
 partitions}

Other Fun: Right

volumes => allows

admin to move these  
 across servers

vol map: <sup>cached</sup> on clients

(vol => machine) relation

(Compare: workloads)

caching:  
 before using file in  
 cache, check if  
 "ok" w/  
 server

measure!  
 => too slow,  
 does not scale

(60% of calls to  
 server -> validate  
 cache  
 entry)

=> load imbalance

=> cpu load  
 (path traversal)

how to fix?

change the protocol

=> too many interactions w/ server

=> interactions are inefficient

AFS v2

same basic structure (whole file caching on disks)

consistency: callbacks

client: when file is accessed first time,  
 obtain callback from server

(+ promise)

=> assume up-to-date unless told  
 otherwise

server: <sup>when</sup> ~~before~~ mod occurs,  
 notify all clients with callback

name resolution:

old: full path => server (too inefficient  
 high load on server (CPU))

new: piece @ time

"name" => FID (vol #, vnode #, unique ID)